

# Cheat Sheet: Fundamentals of Building AI Agents using RAG and LangChain

Package/Method	Description	Code example
Generate text	This code snippet generates text sequences based on the input and doesn't compute the gradient to generate output.	<pre># Generate text output_ids = model.generate(     inputs.input_ids,     attention_mask=inputs.attention_mask,     pad_token_id=tokenizer.eos_token_id,     max_length=50,     num_return_sequences=1 ) output_ids or with torch.no_grad():     outputs = model(**inputs) outputs</pre>
formatting_prompts_func_no_response function	The prompt function generates formatted text prompts from a data set by using the instructions from the data set. It creates strings that include only the instruction and a placeholder for the response.	<pre>def formatting_prompts_func(mydataset):     output_texts = []     for i in range(len(mydataset['instruction'])):         text = (             f"### Instruction:\n{mydataset['instruction'][i]}"             f"\n\n### Response:\n{mydataset['output'][i]}"         )         output_texts.append(text)     return output_texts def formatting_prompts_func_no_response(mydataset):     output_texts = []     for i in range(len(mydataset['instruction'])):         text = (             f"### Instruction:\n{mydataset['instruction'][i]}"             f"\n\n### Response:\n"         )         output_texts.append(text)     return output_texts</pre>
torch.no_grad()	This code snippet helps generate text sequences from the pipeline function. It ensures that the gradient computations are disabled and optimizes the performance and memory usage.	<pre>with torch.no_grad():     # Due to resource limitation, only apply the function on 3 records using "in     pipeline_iterator= gen_pipeline(instructions_torch[:3],                                     max_length=50, # this is set to 50 due to resour                                     num_beams=5,                                     early_stopping=True,)  generated_outputs_lora = [] for text in pipeline_iterator:     generated_outputs_lora.append(text[0]["generated_text"])</pre>
mixtral-8x7b-instruct-v01 watsonx.ai inference model object	Adjusts the parameters to push the limits of creativity and response length.	<pre>model_id = 'mistralai/mixtral-8x7b-instruct-v01' parameters = {     GenParams.MAX_NEW_TOKENS: 256, # this controls the maximum number of tokens     GenParams.TEMPERATURE: 0.5, # this randomness or creativity of the model's r } credentials = {     "url": "https://us-south.ml.cloud.ibm.com" } project_id = "skills-network" model = ModelInference(     model_id=model_id,     params=parameters,     credentials=credentials,     project_id=project_id )</pre>
String prompt templates	Used to format a single string and are generally used for simpler inputs.	<pre>from langchain_core.prompts import PromptTemplate prompt = PromptTemplate.from_template("Tell me one {adjective} joke about {topic}") input_ = {"adjective": "funny", "topic": "cats"} # create a dictionary to store prompt.invoke(input_)</pre>
Chat prompt templates	Used to format a list of messages. These "templates" consist of a list of templates themselves.	<pre>from langchain_core.prompts import ChatPromptTemplate prompt = ChatPromptTemplate.from_messages([     ("system", "You are a helpful assistant"),     ("user", "Tell me a joke about {topic}") ]) input_ = {"topic": "cats"} prompt.invoke(input_)</pre>
Messages place holder	This prompt template is responsible for adding a list of messages in a particular place. But if you want the user to pass in a list of messages that you	<pre>from langchain_core.prompts import MessagesPlaceholder from langchain_core.messages import HumanMessage prompt = ChatPromptTemplate.from_messages([     ("system", "You are a helpful assistant"),     MessagesPlaceholder("msgs") ]) input_ = {"msgs": [HumanMessage(content="What is the day after Tuesday?")]}</pre>



text_splitter	<p>into a larger chunk until you reach a certain size (as measured by some function).</p> <ul style="list-style-type: none"> <li>Once you reach that size, make that chunk its own piece of text and start creating a new chunk with some overlap (to keep context between chunks).</li> </ul>	<pre>text_splitter = CharacterTextSplitter(chunk_size=200, chunk_overlap=20, separator=" ") chunks = text_splitter.split_documents(document) print(len(chunks))</pre>
Embedding models	<p>Embedding models are specifically designed to interface with text embeddings. Embeddings generate a vector representation for a given piece of text. This is advantageous as it allows you to conceptualize text within a vector space. Consequently, you can perform operations such as semantic search, where you identify pieces of text that are most similar within the vector space.</p>	<pre>from ibm_watsonx_ai.metanames import EmbedTextParamsMetaNames embed_params = {     EmbedTextParamsMetaNames.TRUNCATE_INPUT_TOKENS: 3,     EmbedTextParamsMetaNames.RETURN_OPTIONS: {"input_text": True}, } from langchain_ibm import WatsonxEmbeddings watsonx_embedding = WatsonxEmbeddings(     model_id="ibm/slate-125m-english-rtrvr",     url="https://us-south.ml.cloud.ibm.com",     project_id="skills-network",     params=embed_params, )</pre>
Vector store-backed retriever	<p>A retriever that uses a vector store to retrieve documents. It is a lightweight wrapper around the vector store class to make it conform to the retriever interface. It uses the search methods implemented by a vector store, like similarity search and MMR (maximum marginal relevance), to query the texts in the vector store. Since we've constructed a vector store docsearch, it's very easy to construct a retriever.</p>	<pre>retriever = docsearch.as_retriever() docs = retriever.invoke("Langchain")A vector store retriever is a retriever that Since we've constructed a vector store docsearch, it's very easy to construct a</pre>
ChatMessageHistory class	<p>One of the core utility classes underpinning most (if not all) memory modules is the ChatMessageHistory class. This super lightweight wrapper provides convenient methods for saving HumanMessages, AIMessages, and then fetching them all.</p>	<pre>from langchain.memory import ChatMessageHistory chat = mixtral_llm history = ChatMessageHistory() history.add_ai_message("hi!") history.add_user_message("what is the capital of France?")</pre>
langchain.chains	<p>This code snippet uses a LangChain library for building language model applications, creating a chain to generate popular dish recommendations based on the specified locations. It also configures model inference settings for</p>	<pre>from langchain.chains import LLMChain template = """Your job is to come up with a classic dish from the area that the             {location}             YOUR RESPONSE: """ prompt_template = PromptTemplate(template=template, input_variables=['location']) # chain 1 location_chain = LLMChain(llm=mixtral_llm, prompt=prompt_template, output_key='m') location_chain.invoke(input={'location': 'China'})</pre>

	further processing.	
Simple sequential chain	Sequential chains allow the output of one LLM to be used as the input for another. This approach is beneficial for dividing tasks and maintaining the focus of your LLM.	<pre>from langchain.chains import SequentialChain template = """Given a meal {meal}, give a short and simple recipe on how to make YOUR RESPONSE: """ prompt_template = PromptTemplate(template=template, input_variables=['meal']) # chain 2 dish_chain = LLMChain(llm=mixtral_llm, prompt=prompt_template, output_key='recipe') template = """Given the recipe {recipe}, estimate how much time I need to cook it YOUR RESPONSE: """ prompt_template = PromptTemplate(template=template, input_variables=['recipe']) # chain 3 recipe_chain = LLMChain(llm=mixtral_llm, prompt=prompt_template, output_key='time') # overall chain overall_chain = SequentialChain(chains=[location_chain, dish_chain, recipe_chain], input_variables=['location'], output_variables=['meal', 'recipe', 'time'] verbose= True)</pre>
load_summarize_chain	This code snippet uses LangChain library for loading and using a summarization chain with a specific language model and chain type. This chain type will be applied to web data to print a resulting summary.	<pre>from langchain.chains.summarize import load_summarize_chain chain = load_summarize_chain(llm=mixtral_llm, chain_type="stuff", verbose=False) response = chain.invoke(web_data) print(response['output_text'])</pre>
TextClassifier	Represents a simple text classifier that uses an embedding layer, a hidden linear layer with a ReLU activation, and an output linear layer. The constructor takes the following arguments: num_class: The number of classes to classify. freeze: Whether to freeze the embedding layer.	<pre>from torch import nn class TextClassifier(nn.Module):     def __init__(self, num_classes, freeze=False):         super(TextClassifier, self).__init__()         self.embedding = nn.Embedding.from_pretrained(glove_embedding.vectors.to( # An example of adding additional layers: A linear layer and a ReLU acti         self.fcl = nn.Linear(in_features=100, out_features=128)         self.relu = nn.ReLU()         # The output layer that gives the final probabilities for the classes         self.fc2 = nn.Linear(in_features=128, out_features=num_classes)     def forward(self, x):         # Pass the input through the embedding layer         x = self.embedding(x)         # Here you can use a simple mean pooling         x = torch.mean(x, dim=1)         # Pass the pooled embeddings through the additional layers         x = self.fcl(x)         x = self.relu(x)         return self.fc2(x)</pre>
Train the model	This code snippet outlines the function to train a machine learning model using PyTorch. This function trains the model over a specified number of epochs, tracks them, and evaluates the performance on the data set.	<pre>def train_model(model, optimizer, criterion, train_dataloader, valid_dataloader, cum_loss_list = [] acc_epoch = [] best_acc = 0 file_name = model_name for epoch in tqdm(range(1, epochs + 1)):     model.train()     cum_loss = 0     for _, (label, text) in enumerate(train_dataloader):         optimizer.zero_grad()         predicted_label = model(text)         loss = criterion(predicted_label, label)         loss.backward()         torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1)         optimizer.step()         cum_loss += loss.item()     #print("Loss:", cum_loss)     cum_loss_list.append(cum_loss)     acc_val = evaluate(valid_dataloader, model, device)     acc_epoch.append(acc_val)     if acc_val &gt; best_acc:         best_acc = acc_val         print(f"New best accuracy: {acc_val:.4f}")         #torch.save(model.state_dict(), f"{model_name}.pth")     #save_list_to_file(cum_loss_list, f"{model_name}_loss.pkl")     #save_list_to_file(acc_epoch, f"{model_name}_acc.pkl")</pre>
	This code snippet defines function 'llm_model' for generating text using the language model	<pre>def llm_model(prompt_txt, params=None):     model_id = 'mistralai/mixtral-8x7b-instruct-v01'     default_params = {         "max_new_tokens": 256,         "min_new_tokens": 0,         "temperature": 0.5,         "top_p": 0.2,         "top_k": 1     }     if params:         default_params.update(params)     parameters = {         GenParams.MAX_NEW_TOKENS: default_params["max_new_tokens"], # this cont</pre>

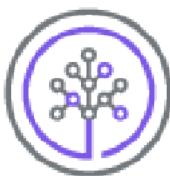
llm_model	from the mistral.ai platform, specifically the 'mistral-8x7b-instruct-v01' model. The function helps in customizing generating parameters and interacts with IBM Watson's machine learning services.	<pre> GenParams.MIN_NEW_TOKENS: default_params["min_new_tokens"], # this controls the minimum number of tokens generated GenParams.TEMPERATURE: default_params["temperature"], # this randomness controls the temperature of the generation GenParams.TOP_P: default_params["top_p"], GenParams.TOP_K: default_params["top_k"] } credentials = {     "url": "https://us-south.ml.cloud.ibm.com" } project_id = "skills-network" model = Model(     model_id=model_id,     params=parameters,     credentials=credentials,     project_id=project_id ) mixtral_llm = WatsonxLLM(model=model) response = mixtral_llm.invoke(prompt_txt) return response </pre>
Zero-shot prompt	Zero-shot learning is crucial for testing a model's ability to apply its pre-trained knowledge to new, unseen tasks without additional training. This capability is valuable for gauging the model's generalization skills.	<pre> prompt = """Classify the following statement as true or false: 'The Eiffel Tower is located in Berlin.' Answer: """ response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n") </pre>
One-shot prompt	One-shot learning example where the model is given a single example to help guide its translation from English to French. The prompt provides a sample translation pairing, "How is the weather today?" translated to "Comment est le temps aujourd'hui?" This example serves as a guide for the model to understand the task context and desired format. The model is then tasked with translating a new sentence, "Where is the nearest supermarket?" without further guidance.	<pre> params = {     "max_new_tokens": 20,     "temperature": 0.1, } prompt = """Here is an example of translating a sentence from English to French: English: "How is the weather today?" French: "Comment est le temps aujourd'hui?" Now, translate the following sentence from English to French: English: "Where is the nearest supermarket?"""" response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n") </pre>
Few-shot prompt	This code snippet classifies emotions using a few-shot learning approach. The prompt includes various examples where statements are associated with their respective emotions.	<pre> #parameters `max_new_tokens` to 10, which constrains the model to generate brief responses params = {     "max_new_tokens": 10, } prompt = """Here are few examples of classifying emotions in statements: Statement: 'I just won my first marathon!' Emotion: Joy Statement: 'I can't believe I lost my keys again.' Emotion: Frustration Statement: 'My best friend is moving to another country.' Emotion: Sadness Now, classify the emotion in the following statement: Statement: 'That movie was so scary I had to cover my eyes.'""" response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n") </pre>
Chain-of-thought (CoT) prompting	The Chain-of-Thought (CoT) prompting technique, designed to guide the model through a sequence of reasoning steps to solve a problem.  The CoT technique involves structuring the prompt by	<pre> params = {     "max_new_tokens": 512,     "temperature": 0.5, } prompt = """Consider the problem: 'A store had 22 apples. They sold 15 apples to a customer. How many apples are there now?' Break down each step of your calculation </pre>

	<p>instructing the model to "Break down each step of your calculation." This encourages the model to include explicit reasoning steps, mimicking human-like problem-solving processes.</p>	<pre>""" response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n")</pre>
Self-consistency	<p>This code snippet determines the consistent result for age-related problems and generates multiple responses. The 'params' dictionary specifies the maximum number of tokens to generate responses.</p>	<pre>params = {     "max_new_tokens": 512, } prompt = """When I was 6, my sister was half of my age. Now I am 70, what age is Provide three independent calculations and explanations, then determ """  response = llm_model(prompt, params) print(f"prompt: {prompt}\n") print(f"response : {response}\n")</pre>
Prompt template	<p>A key concept in LangChain, it helps to translate user input and parameters into instructions for a language model. This can be used to guide a model's response, helping it understand the context and generate relevant and coherent language-based output.</p>	<pre>model_id = 'mistralai/mixtral-8x7b-instruct-v01' parameters = {     GenParams.MAX_NEW_TOKENS: 256, # this controls the maximum number of tokens     GenParams.TEMPERATURE: 0.5, # this randomness or creativity of the model's r } credentials = {     "url": "https://us-south.ml.cloud.ibm.com" } project_id = "skills-network" model = Model(     model_id=model_id,     params=parameters,     credentials=credentials,     project_id=project_id ) mixtral_llm = WatsonxLLM(model=model) mixtral_llm</pre>
Text summarization	<p>Text summarization agent designed to help summarize the content you provide to the LLM. You can store the content to be summarized in a variable, allowing for repeated use of the prompt.</p>	<pre>content = """ The rapid advancement of technology in the 21st century has transformed . Innovations such as artificial intelligence, machine learning, and the I. For instance, AI-powered diagnostic tools are improving the accuracy and Moreover, online learning platforms are making education more accessible These technological developments are not only enhancing productivity but """  template = """Summarize the {content} in one sentence. """  prompt = PromptTemplate.from_template(template) llm_chain = LLMChain(prompt=prompt, llm=mixtral_llm) response = llm_chain.invoke(input = {"content": content}) print(response["text"])</pre>
Question answering	<p>An agent that enables the LLM to learn from the provided content and answer questions based on what it has learned. Occasionally, if the LLM does not have sufficient information, it might generate a speculative answer. To manage this, you'll specifically instruct it to respond with "Unsure about the answer" if it is uncertain about the correct response.</p>	<pre>content = """ The solar system consists of the Sun, eight planets, their moons, dwarf : The inner planets—Mercury, Venus, Earth, and Mars—are rocky and solid. The outer planets—Jupiter, Saturn, Uranus, and Neptune—are much larger a """  question = "Which planets in the solar system are rocky and solid?" template = """ Answer the {question} based on the {content}. Respond "Unsure about answer" if not sure about the answer. Answer: """  prompt = PromptTemplate.from_template(template) output_key = "answer" llm_chain = LLMChain(prompt=prompt, llm=mixtral_llm, output_key=output_key) response = llm_chain.invoke(input = {"question":question , "content": content}) print(response["answer"])</pre>
Code generation	<p>An agent that is designed to generate SQL queries based on given descriptions. It interprets the requirements from your input and translates them into executable SQL code.</p>	<pre>description = """ Retrieve the names and email addresses of all customers from the 'customers' table. The table 'purchases' contains a column 'purchase_date' """  template = """ Generate an SQL query based on the {description} SQL Query: """  prompt = PromptTemplate.from_template(template) output_key = "query" llm_chain = LLMChain(prompt=prompt, llm=mixtral_llm, output_key=output_key) response = llm_chain.invoke(input = {"description":description}) print(response["query"])</pre>

Role playing	Configures the LLM to assume specific roles as defined by us, enabling it to follow predetermined rules and behave like a task-oriented chatbot.	<pre> role = """     game master """ tone = "engaging and immersive" template = """     You are an expert {role}. I have this question {question}. I would l     Answer:  """ prompt = PromptTemplate.from_template(template) output_key = "answer" llm_chain = LLMChain(prompt=prompt, llm=mixtral_llm, output_key=output_key) </pre>
class_names	This code snippet maps numerical labels to their corresponding textual descriptions to classify tasks. This code helps in machine learning to interpret the output model, where the model's predictions are numerical and should be presented in a more human-readable format.	<pre> class_names = {0: "negative", 1: "positive"} class_names </pre>
read_and_split_text	Involves opening the file, reading its contents, and splitting the text into individual paragraphs. Each paragraph represents a section of the company policies. You can also filter out any empty paragraphs to clean your data set.	<pre> def read_and_split_text(filename):     with open(filename, 'r', encoding='utf-8') as file:         text = file.read()     # Split the text into paragraphs (simple split by newline characters)     paragraphs = text.split('\n')     # Filter out any empty paragraphs or undesired entries     paragraphs = [para.strip() for para in paragraphs if len(para.strip()) &gt; 0]     return paragraphs # Read the text file and split it into paragraphs paragraphs = read_and_split_text('companyPolicies.txt') paragraphs[0:10] </pre>
encode_contexts	This code snippet encodes a list of texts into embeddings using content_tokenizer and context_encoder. This code helps iterate through each text in the input list, tokenizes and encodes it, and then appends the pooler_output to the embeddings list. The resulting embeddings get stored in the context_embeddings variables and generate embeddings from text data for various natural language processing (NLP) applications.	<pre> def encode_contexts(text_list):     # Encode a list of texts into embeddings     embeddings = []     for text in text_list:         inputs = context_tokenizer(text, return_tensors='pt', padding=True, truncation=True)         outputs = context_encoder(**inputs)         embeddings.append(outputs.pooler_output)     return torch.cat(embeddings).detach().numpy() # you would now encode these paragraphs to create embeddings. context_embeddings = encode_contexts(paragraphs) </pre>
import faiss	FAISS (Facebook AI Similarity Search) is an efficient library developed by Facebook for similarity search and clustering of dense vectors. FAISS is designed for fast similarity search, which is particularly valuable when dealing with large data sets. It is highly suitable for tasks in natural language processing where retrieval speed is critical. It effectively handles large volumes of data,	<pre> import faiss # Convert list of numpy arrays into a single numpy array embedding_dim = 768 # This should match the dimension of your embeddings context_embeddings_np = np.array(context_embeddings).astype('float32') # Create a FAISS index for the embeddings index = faiss.IndexFlatL2(embedding_dim) index.add(context_embeddings_np) # Add the context embeddings to the index </pre>

	maintaining performance even as data set sizes increase.	
search_relevant_contexts	This code snippet is useful in searching relevant contexts for a given question. It tokenizes the question using the question_tokenizer, encodes the question using question_encoder, and searches an index for retrieving the relevant context based on question embedding.	<pre>def search_relevant_contexts(question, question_tokenizer, question_encoder, index):     """     Searches for the most relevant contexts to a given question.     Returns:     tuple: Distances and indices of the top k relevant contexts.     """     # Tokenize the question     question_inputs = question_tokenizer(question, return_tensors='pt')     # Encode the question to get the embedding     question_embedding = question_encoder(**question_inputs).pooler_output.detach()     # Search the index to retrieve top k relevant contexts     D, I = index.search(question_embedding, k)     return D, I</pre>
generate_answer_without_context	This code snippet generates responses using the entered prompt without requiring additional context. It tokenizes the input questions using the tokenizer, generates the output text using the model, and decodes the generated text to obtain the answer.	<pre>def generate_answer_without_context(question):     # Tokenize the input question     inputs = tokenizer(question, return_tensors='pt', max_length=1024, truncation=True)     # Generate output directly from the question without additional context     summary_ids = model.generate(inputs['input_ids'], max_length=150, min_length=10)     # Decode and return the generated text     answer = tokenizer.decode(summary_ids[0], skip_special_tokens=True)     return answer</pre>
Generating answers with DPR contexts	Answers are generated when the model utilizes contexts retrieved via DPR, which are expected to enhance the answer's relevance and depth:	<pre>def generate_answer(contexts):     # Concatenate the retrieved contexts to form the input to BART     input_text = ' '.join(contexts)     inputs = tokenizer(input_text, return_tensors='pt', max_length=1024, truncation=True)     # Generate output using BART     summary_ids = model.generate(inputs['input_ids'], max_length=150, min_length=10)     return tokenizer.decode(summary_ids[0], skip_special_tokens=True)</pre>
aggregate_embeddings function	The function aggregate_embeddings takes token indices and their corresponding attention masks, and uses a BERT model to convert these tokens into word embeddings. It then filters out the embeddings for zero-padded tokens and computes the mean embedding for each sequence. This helps in reducing the dimensionality of the data while retaining the most important information from the embeddings.	<pre>def aggregate_embeddings(input_ids, attention_masks, bert_model=bert_model):     """     Converts token indices and masks to word embeddings, filters out zero-padded     and aggregates them by computing the mean embedding for each input sequence.     """     mean_embeddings = []     # Process each sequence in the batch     print('number of inputs', len(input_ids))     for input_id, mask in tqdm(zip(input_ids, attention_masks)):         input_ids_tensor = torch.tensor([input_id]).to(DEVICE)         mask_tensor = torch.tensor([mask]).to(DEVICE)         with torch.no_grad():             # Obtain the word embeddings from the BERT model             word_embeddings = bert_model(input_ids_tensor, attention_mask=mask_tensor)             # Filter out the embeddings at positions where the mask is zero             valid_embeddings_mask = mask_tensor[0] != 0             valid_embeddings = word_embeddings[valid_embeddings_mask, :]             # Compute the mean of the filtered embeddings             mean_embedding = valid_embeddings.mean(dim=0)                                mean_embeddings.append(mean_embedding)     # Concatenate the mean embeddings from all sequences in the batch     aggregated_mean_embeddings = torch.cat(mean_embeddings)     return aggregated_mean_embeddings</pre>
text_to_emb	Designed to convert a list of text strings into their corresponding embeddings using a pre-defined tokenizer.	<pre>def text_to_emb(list_of_text, max_input=512):     data_token_index = tokenizer.batch_encode_plus(list_of_text, add_special_tokens=False)     return data_token_index</pre>
process_song	Convert both the predefined appropriateness questions and the song lyrics into "RAG embeddings" and measure the similarity between them to determine the appropriateness.	<pre>import re def process_song(song):     # Remove line breaks from the song     song_new = re.sub(r'\n', ' ', song)     # Remove single quotes from the song     song_new = [song_new.replace("'", "")]</pre>

RAG_QA	This code snippet performs question-answering using question embeddings and provides embeddings. It helps reshape the results for processing, sorting the indices in descending order, and printing the top 'n-responses' based on the highest dot product values.	<pre>def RAG_QA(embeddings_questions, embeddings, n_responses=3):     # Calculate the dot product between the question embeddings and the provided     dot_product = embeddings_questions @ embeddings.T     # Reshape the dot product results to a 1D tensor for easier processing.     dot_product = dot_product.reshape(-1)     # Sort the indices of the dot product results in descending order (setting d     sorted_indices = torch.argsort(dot_product, descending=True)     # Convert sorted indices to a list for easier iteration.     sorted_indices = sorted_indices.tolist()     # Print the top 'n_responses' responses from the sorted list, which correspo     for index in sorted_indices[:n_responses]:         print(yes_responses[index])</pre>
model_name_or_path	This code snippet defines the model name to 'gpt2' and initializes the token and model using the GPT-2 model. In this code, add special tokens for padding by keeping the maximum sequence length to 1024.	<pre># Define the model name or path model_name_or_path = "gpt2" # Initialize tokenizer and model tokenizer = GPT2Tokenizer.from_pretrained(model_name_or_path, use_fast=True) model = GPT2ForSequenceClassification.from_pretrained(model_name_or_path, num_la # Add special tokens if necessary tokenizer.pad_token = tokenizer.eos_token model.config.pad_token_id = model.config.eos_token_id # Define the maximum length max_length = 1024</pre>
add_combined_columns	This code snippet combines the prompt with chosen and rejected responses in a data set example. It combines with the 'Human:' and 'Assistant:' for clarity. This function modifies each example in the 'train' split the data set by creating new columns 'prompt_chosen' and 'prompt_rejected' with the combined text.	<pre># Define a function to combine 'prompt' with 'chosen' and 'rejected' responses def add_combined_columns(example):     # Combine 'prompt' with 'chosen' response, formatting it with "Human:" and ".     example['prompt_chosen'] = "\n\nHuman: " + example["prompt"] + "\n\nAssistan     # Combine 'prompt' with 'rejected' response, formatting it with "Human:" and "     example['prompt_rejected'] = "\n\nHuman: " + example["prompt"] + "\n\nAssist     # Return the modified example     return example # Apply the function to each example in the 'train' split of the dataset dataset['train'] = dataset['train'].map(add_combined_columns)</pre>
RetrievalQA	This code snippet creates an example for 'RetrievalQA' using a language model and document retriever.	<pre>qa = RetrievalQA.from_chain_type(llm=flan_llm,                                     chain_type="stuff",                                     query = "what is mobile policy?"                                     qa.invoke(query)</pre> <span style="float: right;">ret</span>



# Skills Network